

Significant Error: Citations to Research With Publicized Statistical Errors

Ken Cor Gaurav Sood

October 2026

Abstract

Publicizing a statistical error could sharply reduce reliance on affected papers or prompt researchers to qualify their citations. We examine both responses after a 2011 critique of a common mistake in neuroscience: inferring that two effects differ because one is statistically significant and the other is not. Papers flagged for the mistake continued to receive citations well beyond a handful of highly cited cases. Their median annual citations rose from 5 in 2010 to 13–17 during 2012–2015; comparison papers rose from 4 to 10–11. In a separate sample of citations to flagged papers, 94 of 95 valid completed ratings recorded no acknowledgment of concerns. These findings document widespread continued use with little recorded qualification. Comparisons with unflagged papers leave the size of any citation penalty uncertain; continued citation does not establish that publicity had no effect. We also examine 67 political-science papers assessed in an audit of instrumental-variable methods. Papers flagged for statistical problems continued to receive citations after the audit appeared, but the comparisons leave uncertain how much, if at all, the audit reduced their citation growth.

We thank Danielle Portnoy and Xiaoran Huang for research assistance, and Andrew Gelman, Kabir Khanna, and Daniel Stone for comments. Replication materials: https://github.com/recite/sig_fault.

1 Introduction

Publicizing a statistical error gives researchers a reason to reconsider their reliance on affected findings. A substantial response could take two forms: fewer citations to papers containing the error, or more citations that acknowledge the problem. We examine both after a prominent methodological critique in neuroscience.

The mistake is to treat a statistically significant effect in one group and a nonsignificant effect in another as evidence that the groups differ. That conclusion requires a test of the difference itself. Significance in one group can reflect greater precision rather than a larger effect (Gelman and Stern, 2006). The missing test leaves the claimed contrast unsupported by that argument; it does not establish that every conclusion in the paper is false.

Nieuwenhuis et al. (2011) identified 157 papers in which a comparison of effects was relevant and classified 79 as making the mistake. They supplied us with the paper-level classifications. Researchers who understand the critique can recognize the mistake in an original paper and reconsider its use. The critique therefore need not name each affected paper to influence citations.

The affected papers remained widely cited, and recorded qualifications were rare. Median as well as mean citations increased after the critique; 69 of 76 flagged papers received more citations per year in 2012–2015 than in 2010. Almost all valid completed ratings in our citation-context sample recorded no concern. The comparison with unflagged papers is less conclusive about how much publicity changed citation growth. We distinguish that causal question from the observed persistence of largely unqualified citation.

2 Data and comparison

The source review covers behavioral, systems, and cognitive neuroscience papers published in 2009–2010 in *Nature*, *Science*, *Nature Neuroscience*, *Neuron*, and sampled issues of the *Journal of Neuroscience*. The classification concerns the particular statistical inference reviewed, not the validity of every result in a paper.

The archived citation exports cover 155 of the 157 classified papers. Two further histories contain the intended cited paper itself and records predating it, indicating invalid citation searches. Excluding these histories leaves 76 flagged and 77 comparison papers. We count distinct citing documents per paper and year, retain citation-free years within covered histories, and exclude the incomplete 2016 export. Appendix A documents coverage, duplicates, and linkage checks.

The critique appeared online on August 26, 2011. We use 2010 as the baseline and 2012–2015 as the post-publicity period, omitting the transition year 2011. The main regression is Poisson pseudo-maximum likelihood with article and calendar-year fixed effects. It compares proportional citation changes in flagged and comparison papers; uncertainty is clustered by article. A relative post/pre citation ratio below one indicates slower proportional growth for flagged papers. Appendix B gives the specification and alternative comparisons.

Interpreting that contrast as an effect of publicity requires comparable proportional citation trajectories without the critique. Error status was not randomized, and papers may differ in topic,

prominence, and ordinary citation growth. A general warning could also affect both groups; the comparison concerns their differential response, not the critique’s total influence on science.

3 Continued citation and qualification

3.1 Citation trajectories

Continued citation extends beyond a few prominent papers. Flagged papers’ median annual citations rose from 5 in 2010 to between 13 and 17 in each year from 2012 through 2015 (Figure 1). Comparison papers’ medians rose from 4 to 10–11. Mean citations also rose: from 9.4 to a post-period annual average of 21.4 for flagged papers, and from 6.5 to 16.7 for comparison papers. For 91% of flagged papers, the post-period annual average exceeded the baseline.




Figure 1: Median and mean annual citations to 76 flagged and 77 comparison papers, including citation-free years. The vertical line marks August 2011, when the critique appeared. The 2009 series includes prepublication zeros for papers published in 2010; the regression baseline is 2010. Lines describe the covered papers rather than a probability sample of all research.

The main regression estimates a 11.5% smaller post/pre citation ratio for flagged papers (95% interval [-31.6, 14.4]%). Thus, widespread continued citation coexists with considerable uncertainty about a relative citation penalty. The medians establish that continued use was typical, while the regression addresses how growth compares with other papers. Neither an observed increase nor an imprecise comparison establishes that publicity had no effect.

Allowing citation trajectories to differ by journal and publication cohort produces estimates from -11.4% to -7.7%. Absolute citation gains instead favor flagged papers by 1.8 citations per paper per year (95% interval [-2.1, 5.7]): flagged papers started from a higher citation level, so absolute and proportional changes answer different questions. The appendix reports the full comparisons, later response windows, and bounds on the magnitude of a relative decline.

3.2 Qualifications in citing papers

The archived citation-context sample contains 100 relationships selected after 2011, including 7 from the incomplete 2016 export. All concern citations to flagged papers. Of 95 valid completed ratings, 94 recorded no acknowledgment of concerns and 1 recorded acknowledgment. Unqualified citation was therefore common in the observed sample. The protocol records any concern about the cited paper, not only the specific interaction error.

A citation without criticism need not endorse the faulty inference. It may use another finding, a method, or background information. Nor do the codes reveal whether the authors recognized the mistake. Because the coding sample covers only the post-publicity period, it establishes the rarity of recorded qualifications after publicity, not how their frequency changed. Appendix A reports missing ratings and uncertainty.

4 Additional evidence from an instrumental-variables audit

A second methodological critique permits a comparison across fields and statistical problems. Lal et al. (2024) reanalyze instrumental-variable designs in 67 political-science papers. We recover citation histories for that entire assessed-paper cohort. These classifications concern instrument strength and sensitivity of inference; they do not establish that the original substantive conclusions are false. The comparison therefore extends the citation design to adverse diagnostics while keeping that distinction from demonstrated errors explicit.

We examine two main classifications. The first flags an effective first-stage F below 10, a screening threshold rather than a universal criterion for weak instruments. The second flags a reported estimate that is significant at $p < .05$ under the archived analytic test but not under the archived Anderson–Rubin or applicable tF procedure. This second comparison is restricted to papers with at least one analytically significant estimate. A paper satisfies a diagnostic if any of its assessed designs does. Table 1 also reports their union and the narrower Anderson–Rubin comparison; these overlapping definitions are alternative analyses of one audit.

The audit appeared online in May 2024, but an earlier manuscript is dated July 2021. We compare 2023 with 2025 and omit 2024. This contrast concerns citation changes around formal publication, not the effect of first exposure to previously unknown criticism. We count OpenAlex articles and reviews, deduplicate within original-paper citation histories, retain observed zero years, and fit the same article/year fixed-effects Poisson specification with article-clustered uncertainty. All included papers were published before the baseline year. Appendix C describes counting and diagnostic checks.

Papers flagged by the audit continued to receive citations. Under the two main definitions, their mean annual citations rose from 7.0 to 7.2 and from 7.1 to 10.3 between 2023 and 2025; their medians changed little. The comparisons with other assessed papers leave uncertain how much, if at all, the audit reduced citation growth. Both main intervals span reductions and increases, and the estimates change substantially when measured from an earlier year (Appendix C). The opposite signs in Table 1 do not establish that researchers responded differently to different statistical problems.

Table 1: Citation changes around the 2024 instrumental-variables audit

Diagnostic	Flagged	Comparison	Change (%)	95% interval
Effective F below 10	8	59	-14.5	[-36.0, 14.2]
Inferential sensitivity	11	44	23.9	[-7.1, 65.3]
Either diagnostic	17	50	11.8	[-14.6, 46.3]
AR-only sensitivity	3	52	-44.0	[-62.6, -16.0]

Change is $100(e^{\hat{\beta}} - 1)$, comparing flagged and comparison papers’ 2025/2023 citation ratios. Intervals use article-clustered covariance and t_{G-1} critical values. The inferential-sensitivity and AR-only rows restrict the comparison population to papers with an analytically significant estimate. The rows overlap and are not independent studies. AR-only inference is especially fragile because only 3 papers meet that diagnostic.

The timing also differs from the neuroscience study: the IV audit names the papers it assesses and circulated before formal publication. The historical neuroscience exports and the IV citation histories also come from different databases. A full-cohort comparison with OpenCitations yields almost the same neuroscience growth contrast as the historical exports (Appendix C); the OpenAlex comparison remains incomplete.

5 Individually publicized errors

The Institute for Replication extension currently supports matched comparisons for only 3 affected papers. This is a pilot, too small and selected to establish how citations generally respond to publicized errors. The source inventory is much larger, but most records still lack a verified consequential error, an established first-disclosure date, sufficient citation history, or complete controls. Appendix E retains the case comparisons and synthetic-control checks.

6 Discussion

The affected papers continued to circulate after their statistical mistake was publicized. This pattern appears in median as well as mean citations and is accompanied by very little recorded qualification in the context sample. The evidence is inconsistent with wholesale abandonment of these papers. It is less informative about how many citations publicity prevented.

Citation growth already differed before the critique. Among papers published in 2009, flagged papers’ increase from 2009 to 2010 exceeded that of comparison papers by 4.3 citations (95% interval [0.1, 8.6]). That short publication-year baseline cannot establish the later counterfactual trajectory. Awareness is also unobserved, and some later citations may concern findings unaffected by the mistake. These limits leave room for a citation penalty even as the affected papers remain widely used.

The practical question is how methodological criticism reaches the point of use. Readers can recognize an error from the original analysis, but doing so requires connecting a general warning to a particular paper. Our results document continued citation with little recorded acknowledgment; they do not identify which step in that process accounts for it.

References

- Andrew Gelman and Hal Stern. The difference between “significant” and “not significant” is not itself statistically significant. *The American Statistician*, 60(4):328–331, 2006. doi: 10.1198/000313006X152649.
- Apoorva Lal, Mackenzie Lockhart, Yiqing Xu, and Ziwen Zu. How much should we trust instrumental variable estimates in political science? practical advice based on 67 replicated studies. *Political Analysis*, 32:521–540, 2024. doi: 10.1017/pan.2024.2.
- Sander Nieuwenhuis, Birte U. Forstmann, and Eric-Jan Wagenmakers. Erroneous analyses of interactions in neuroscience: a problem of significance. *Nature Neuroscience*, 14(9):1105–1107, 2011. doi: 10.1038/nn.2886.
- J. M. C. Santos Silva and Silvana Tenreiro. The log of gravity. *The Review of Economics and Statistics*, 88(4):641–658, 2006. doi: 10.1162/rest.88.4.641.

A Data construction and citation-context records

The supplied classification spreadsheet identifies the articles by journal, publication year, page, and study characteristics. The identified Excel and CSV versions agree on the paper-level YES/NO codes. The citation import reads 16,437 database records across the covered sheets, repairs 1 malformed record, and marks 15 duplicate DOI records within cited papers. It removes 3 citation relationships identified as false in the historical context coding. Full source hashes and record-level indicators are generated by the replication scripts.

The unreliable histories are source IDs 23 and 25. Matching the intended papers’ journal, year, and supplied page locates the target papers themselves in their supposed citing-article exports. Those histories also contain 54 records before the target publication year. This combination is consistent with exporting literature-search results instead of the intended citation results. The saved workbook contains no query log with which to reconstruct the precise operation. Deleting only the early records would not establish the validity or completeness of the remaining links. The principal estimates exclude both histories, while the appendix reports estimates retaining them. The other sheets contain neither a target-self match nor a record before the target year under these checks.

Missing sheets, source IDs 95 and 124, remain missing. In covered histories the completed 2010–2015 panel includes 29 paper-years with zero recorded citations. The scripts check that every paper has each required year and that joining the classification does not multiply or discard citation records.

Table 2: Disposition of the historical citation-context sample

Recorded status	Citation relationships
Article unavailable	1
Concern recorded	1
False citation link	3
No concern recorded	94
Uncoded	1

The sample consists of citation relationships from 2012 through the partial 2016 export. Every relationship matches a record in the source workbooks. “No concern recorded” and “Concern recorded” reproduce completed historical codes for valid citation relationships. The other categories preserve missingness and known false links.

The original coding protocol instructed the rater to assess whether a citing paper noted concerns about the cited paper, using the citing text and original article. The distributed file contains the rating and selected contextual text, not the full highlighted PDFs. We retain those ratings rather than claiming to have independently recoded the articles. The available materials contain one set of item-level ratings, so intercoder agreement cannot be independently checked.

Among valid completed ratings, the acknowledgment share is 1.1% (95% Wilson interval [0.2, 5.7]%). Assigning both unresolved, non-false records to either category gives a share between 1.0% and 3.1% among all non-false sampled relationships. These are missing-rating bounds, not sampling intervals. The Wilson interval is a binomial reference interval for sampled citation relationships; repeated source papers, selection, and coding errors are not represented by it.

B Models, sensitivity, and uncertainty

B.1 Proportional changes

For paper i in year t , the main model is

$$E[C_{it} \mid i, t, F_i] = \exp\{\alpha_i + \lambda_t + \beta F_i P_t\},$$

where F_i denotes a flagged paper and P_t denotes 2012–2015. Article effects absorb stable differences in citation levels and year effects absorb common calendar-year changes. Poisson pseudo-maximum likelihood retains zero counts and does not require equality of conditional mean and variance for the robust inference used here (Santos Silva and Tenreiro, 2006). The reported percentage is $100(e^\beta - 1)$, a relative post/pre citation change rather than a percentage change in the number of additional citations gained. In this balanced baseline specification, $e^{\hat{\beta}}$ equals the ratio of the groups’ post/pre mean citation ratios.

The covariance clusters score contributions by paper. We apply the $G/(G-1)$ and $(N-1)/(N-K)$ corrections, counting fixed effects not nested within article clusters in K , and use t_{G-1} critical values. Table 3 reports both two-sided 95% intervals and one-sided 95% lower bounds, transformed from the log scale. These are conditional model-based comparisons, not randomized-treatment intervals. No papers are excluded from the main fit for zero outcomes.

Table 3: Proportional citation comparisons

Specification	Flagged / other	Change (%)	95% interval	Lower bound
Article and year effects	76 / 77	-11.5	[-31.6, 14.4]	-28.7
Journal-by-year effects	76 / 77	-11.4	[-31.4, 14.5]	-28.5
Cohort-by-year effects	76 / 77	-7.7	[-24.4, 12.7]	-22.0
Journal/year and cohort/year effects	76 / 77	-8.2	[-24.8, 12.1]	-22.3
Retain questionable histories	78 / 77	-12.2	[-31.9, 13.2]	-29.1
Retain duplicates	76 / 77	-11.6	[-31.6, 14.3]	-28.7
Retain false links	76 / 77	-11.5	[-31.6, 14.4]	-28.7
2009 publication cohort	44 / 47	-7.2	[-25.0, 14.9]	-22.3
2010 publication cohort	32 / 30	-10.4	[-49.8, 60.2]	-44.8
Post period: 2013–2015	76 / 77	-10.7	[-31.5, 16.4]	-28.5
Post period: 2014–2015	76 / 77	-11.7	[-32.7, 16.0]	-29.7
Potentially serious errors	43 / 77	12.8	[-18.3, 55.6]	-13.9

Change and bounds are percentages, $100(e^\beta - 1)$. The last column is the one-sided 95% lower bound. Every model includes article effects. The first row includes year effects; the next three replace common year effects with journal-by-year effects, publication-cohort-by-year effects, or both. Other rows use the first row’s specification with the stated inclusion or post-period change. The baseline is always 2010. The potentially-serious subgroup retains the same comparison papers and excludes other flagged papers. Standard errors cluster by article.

The main lower bound excludes proportional reductions greater than 28.7% at the one-sided 95% confidence level under this model. This leaves substantial reductions compatible with the data. The journal/cohort-adjusted lower bounds permit reductions of 22.0–28.5%; these bounds do not resolve uncertainty about the counterfactual trend. A stratified bootstrap resampling whole papers within each group gives a two-sided percentile interval of [-31.2, 14.5]% and a one-sided bound

permitting a reduction of 28.3%. It uses 9,999 draws with seed 31415; no draws are discarded. The baseline ratio identity allows each bootstrap estimate to be calculated directly from the resampled paper-level before and after means.

Year-specific comparisons use 2010 as the baseline (Figure 2). A paper with zero citations in both the baseline and the comparison year has no finite Poisson article intercept and contributes no information to the slope; such papers are excluded from that fit and listed in the generated results. The 2011 contrast describes the transition year. The available history is too short to justify fitting article-specific trends from predominantly post-publicity observations.




Figure 2: Year-specific proportional citation changes for flagged relative to comparison papers, using 2010 as the baseline. Points are $100(e^{\beta} - 1)$ from Poisson article/year fixed-effects models; bars are simultaneous 95% Bonferroni intervals across the five contrasts, with article-clustered uncertainty. Papers with zero counts in both compared years contribute no identifying information and are omitted from the corresponding fit. The 2011 contrast is transitional and is excluded from the main post period.

B.2 Absolute changes

For an absolute comparison, each paper’s change is its mean annual citations in 2012–2015 minus its 2010 citations. Comparing the two groups’ mean changes gives 1.8 (95% Welch interval $[-2.1, 5.7]$), allowing unequal variances of paper-level changes. OLS with article and year fixed effects on the same balanced panel reproduces this point estimate exactly; it is another implementation of the same comparison, not independent corroboration. The generated results also report its article-clustered interval. Journal and publication-year adjustment in the paper-level change regression uses HC3 uncertainty.

Removing each source paper in turn gives absolute point estimates from 1.3 to 2.5. A whole-paper stratified bootstrap gives a percentile interval of $[-2.0, 5.6]$. Medians describe typical citation

Table 4: Absolute citation changes under alternative comparisons

Comparison	Flagged / other	Difference	95% interval
Main comparison	76 / 77	1.8	[-2.1, 5.7]
Journal and publication-year adjustment	76 / 77	0.8	[-2.8, 4.4]
Retain both questionable citation histories	78 / 77	1.6	[-2.3, 5.5]
Retain duplicate database records	76 / 77	1.8	[-2.2, 5.7]
Retain known false citation links	76 / 77	1.8	[-2.1, 5.7]
Papers published in 2009	44 / 47	2.5	[-2.3, 7.4]
Papers published in 2010	32 / 30	0.3	[-6.4, 7.0]
Later window: 2013-2015	76 / 77	2.0	[-2.2, 6.2]
Later window: 2014-2015	76 / 77	1.7	[-2.5, 6.0]
Potentially serious errors only	43 / 77	4.3	[-0.1, 8.7]
Pre-critique change: 2009 cohort, 2009 to 2010	44 / 47	4.3	[0.1, 8.6]

Differences are flagged minus comparison, in citations per paper per year. Except for the final row, the baseline is 2010; the post period is 2012–2015 unless otherwise stated. Intervals are pointwise 95% Welch intervals, except the journal/year-adjusted estimate, which uses OLS with HC3 standard errors and a residual-degrees-of-freedom t critical value. The final row compares 2010 with 2009 among papers published in 2009; it is a diagnostic, not a publicity-effect estimate.

levels; neither medians nor these sensitivity checks make error status or publicity exposure random. All analytical choices are retrospective.

Of the distinct citing documents in the main papers’ 2010–2015 records, 13.7% cite more than one study paper. Shared citing documents and related topics can induce dependence across papers. Article clustering and the whole-paper bootstrap do not model those connections. The source review is also a selected set of journals and years, so these comparisons are not probability-sample estimates for all scientific publishing.

C Additional-cohort construction and comparability

The IV replication archive contains multiple designs for some papers. We retain the design records but assign each paper one citation history and equal weight within its diagnostic group. The repository reproduces the point estimates for the 10 designs selected in the separate citation-context pilot. For the Alt design, the instrument is an eight-level factor with seven excluded indicators. Recomputing the joint Anderson–Rubin test with that encoding corrects the archived calculation; its loss-of-significance classification remains the same. The single-instrument tF calculation is inapplicable and does not enter that design’s sensitivity classification. These checks do not constitute independent replication of every diagnostic in the full audit.

The relative citation-growth estimates are -14.5% for the weak-instrument screen and 23.9% for inferential sensitivity. On a fixed older cohort, moving the baseline from 2023 to 2022 changes the weak-instrument estimate from -11.5% to -0.4%, and the sensitivity estimate from 15.7% to -17.0%. Holding the cohort fixed shows that changing the starting year alone can substantially change the comparison.

Document types materially affect the IV comparison. The original broad outcome included preprints, proceedings and book chapters; some books contributed links from numerous chapters

and front matter to the same original paper. After observing this problem and the alternative estimates, we selected articles and reviews as the primary outcome. For the weak-instrument screen, the broader count gives 39.3%, compared with -14.5% for articles and reviews. This was an outcome-aware measurement revision. The repository retains the broader results and a partial sensitivity that collapses identified Oxford book DOI families; neither procedure resolves all cross-version duplication.

The source-comparison protocol preserves the original neuroscience analysis cohort and compares identical papers and years before considering a pooled result. Exact DOI matches identify shared citing relationships, with each database’s publication year retained. Historical exports lack the detailed document-type field needed to impose an identical article/review filter. Missing histories stay missing; a relationship absent from one database is not automatically a false citation. Source discrepancies therefore measure differences between the recorded citation histories rather than bias relative to an assumed ground truth.

The currently available paired subset contains 10 flagged and 0 comparison papers. Among 219 shared citation relationships with conflicting years, OpenAlex dates 218 earlier; publisher-deposited online and print dates account for 97 of these differences. Moving shared eligible citations between years accounts for 15 of the 21 additional baseline-year citations recorded by OpenAlex. This preliminary comparison shows why citation dating can change measured growth. Estimating how replacing the historical counts with OpenAlex changes the flagged-versus-comparison contrast requires paired OpenAlex histories for both groups.

For the 10 paired papers checked against OpenCitations, that index also records 118 of the 149 OpenAlex-only links eligible for the primary article/review count. Across all document types in the comparison frame, it records 194 of 292 OpenAlex-only links. This corroborates many additional records, but shared upstream sources mean agreement across indexes is not independent verification of the citing bibliography. Undated OpenCitations links remain undated, and this check does not replace either source’s annual counts.

A separate OpenCitations comparison covers all 153 historical papers. Its estimated relative-growth contrast is -11.9%, compared with -11.5% in the historical exports. OpenCitations records more citations in both periods, leaving this aggregate comparison almost unchanged. The outcome counts dated citing works across document types; 310 undated relationships remain unassigned. A paired, paper-level bootstrap gives a source-induced change in the estimated growth ratio of -0.4% (95% interval [-8.0, 9.2]%). These intervals condition on the observed cohort and database records. Similar aggregate estimates do not verify individual links or establish either source as ground truth.

For comparison across audits, the repository also reports the contrast between the calendar year before and the calendar year after the warning year, excluding the warning year itself. This uses 2010 and 2012 for neuroscience and 2023 and 2025 for the IV audit. The neuroscience papers published in 2010 have a partial publication-year baseline, so a separate 2009-publication cohort has a complete baseline year. The synthesis below uses these common windows; alternative IV diagnostics are not counted as independent replications.

D Provisional synthesis across audits

The completed cohorts support a provisional summary of their relative citation-growth contrasts. We give each audit equal weight on the log scale: $\bar{\beta} = (\hat{\beta}_N + \hat{\beta}_L)/2$. The reported quantity, $100(e^{\bar{\beta}} - 1)$, transforms the geometric average of their relative growth ratios. It is neither the citation ratio obtained by pooling papers nor the arithmetic average percentage change. Each analysis includes one contrast from each audit; changing the IV diagnostic changes the estimand.

Table 5: Equal-audit synthesis under the two primary IV definitions

IV diagnostic	Neuroscience	IV audit	Combined	95% interval
Effective F below 10	-14.1	-14.5	-14.3	[-29.0, 3.4]
Inferential sensitivity	-14.1	23.9	3.1	[-14.5, 24.4]

All entries are percentage relative growth contrasts. Component estimates compare 2010 with 2012 in neuroscience and 2023 with 2025 in the IV audit. Combined estimates give the two log contrasts equal weight. The intervals condition on these audits and assume independent audit estimates; they do not include uncertainty from choosing audits, databases, timing or diagnostic definitions.

The combined contrast is -14.3% (95% interval [-29.0, 3.4]%) under the instrument-strength screen and 3.1% ([-14.5, 24.4]%) under inferential sensitivity. The difference illustrates why the diagnostic definition matters. The repository also reports the union and AR-only alternatives and repeats the synthesis using only neuroscience papers published in 2009. The AR-only comparison remains exploratory because its IV component contains only 3 flagged papers.

With component variances s_N^2 and s_L^2 , the variance of the equal-audit mean is $(s_N^2 + s_L^2)/4$. Intervals use a Welch–Satterthwaite approximation, with degrees of freedom $(s_N^2 + s_L^2)^2 / (s_N^4/d_N + s_L^4/d_L)$ and component article-cluster degrees of freedom d_N, d_L . This approximation uses cluster-robust component variances and assumes independence across audits. The interval describes estimation uncertainty for these fixed cases; it is not a random-effects prediction interval for other critiques. Different citation sources, incomplete type harmonization, prior circulation of the IV critique, and the distinction between a demonstrated error and an adverse diagnostic remain embedded in the combined estimates. We therefore interpret them as descriptive summaries, not a general causal effect of publicizing errors.

A retrospective sensitivity adds the 3 matched I4R cases described in Appendix E. Its component compares the post/pre ratio of affected papers’ mean citations with that of their weighted controls, giving 22.3%. This is a ratio of group means, not the equal-case average log ratio, which gives 3.9%. The preceding-year relative-growth contrast is +25.5%; omitting the inventor-clusters case changes the post-disclosure contrast to -25.6%. These diagnostics limit its interpretation as a publicity response.

Giving each of the three components one third of the log-scale weight yields -3.5% ([-28.8, 30.7]%) under the IV strength diagnostic and 9.2% ([-19.4, 47.9]%) under inferential sensitivity. This is an equal-component descriptive summary: I4R contains multiple disclosures rather than one audit and uses all-type citation totals. Its variance uses a case-level delta calculation with 3 disclosures; the combined intervals use the same Welch approximation. No original or control DOI is shared across components, but the I4R and IV calendar windows overlap. The intervals assume

zero cross-component covariance and do not resolve the differences in outcomes, exposure clocks or sample selection.

E Matched cases with individually publicized errors

The I4R registry separates catalog listings, article-specific assessments and verified error disclosures. Unavailable reports and unnamed reviewer teams remain unresolved records; overview summaries alone do not establish material errors. Verified cases link the original article to the error, affected claim, source passage, replies and earliest substantiated public disclosure. Material corrections can strengthen or weaken a finding; these are not necessarily refuted papers. A later I4R listing does not reset the disclosure clock. The repository’s claim assessments and disclosure ledger provide the case-level evidence.

The secondary analysis uses cached OpenAlex work records with verified identities, source hashes, retrieval dates and supported annual coverage. Missing histories remain missing. Controls exclude known assessed papers, retracted papers and verified nonresearch items. Matching waits for a complete eligible journal pool and uses up to three controls, selected by pre-disclosure citation levels, growth and text distance with fixed calipers. The matching rules and measurement choice were recorded before these contrasts were estimated; this remains a retrospective analysis. Post-disclosure outcomes do not enter matching.

Table 6: Citation changes after individually publicized errors

Paper	Years	Affected	Controls	Contrast
Parental leave	2020–2022	6 → 9	9.7 → 15.3	-2.7
Inventor clusters	2023–2025	48 → 102	51.3 → 68.0	+37.3
Electrification	2020–2022	34 → 28	26.0 → 29.0	-9.0

Annual citing works of all document types. The years column gives the calendar year before and after first documented disclosure; the disclosure year is excluded. Controls are equally weighted within each matched set. The final column subtracts the control change from the affected paper’s change. These are individual observed contrasts, not independently identified causal effects.

Mean citations rose from 29.3 to 46.3 among affected papers and from 29 to 37.4 among controls. The equal-paper average difference in changes is 8.6 citations (exploratory 95% interval [-56.9, 74]). First differences with matched-set effects reproduce this average. The covariance clusters by original article and disclosure event, retaining original identities when controls recur; the critical value uses the smaller cluster count minus one as degrees of freedom. Inference conditions on selected matches and is fragile with only 3 disclosures.

The preceding year shows a matched growth difference of +5.6 citations, despite matching on pre-period outcomes. Figure 3 shows the differing trajectories. Leaving out the inventor-clusters case changes the average post-disclosure contrast to -5.8. The repository reports all leave-one-disclosure-out contrasts, alternative baselines and longer horizons on common cohorts. These diagnostics prevent treating the positive pooled contrast as evidence that publicity increases citations. They also show why continued citation alone cannot measure the citations that publicity prevented.




Figure 3: Citation paths for the supported I4R cases. Solid lines show each affected paper; dashed lines show its equally weighted matched controls. Counts include citing works of all types. Year zero marks the calendar year of first documented public disclosure, not its exact date. All panels use the same scale. These descriptive counts have no sampling intervals and do not measure endorsement of the erroneous claim.

E.1 Synthetic-control sensitivity

We also fit nonnegative donor weights summing to one to match each affected paper’s pre-disclosure citation path. Donors retain the journal, publication-age, identity and exposure restrictions, but the citation caliper and text ranking are removed. We require at least three full pre-years, first withholding the last year to check prediction and then refitting on all pre-years. This supports 2 cases; inventor clusters has only two full pre-years. A fixed small ridge penalty resolves the many possible donor combinations. All choices are retrospective.

Table 7: Synthetic-control fit and prediction checks

Paper	Pre-fit RMSE	Held-out: actual / predicted	Post: actual / synthetic	Gap
Parental leave	0.0	6.0 / 8.4	9.0 / 9.6	-0.6
Electrification	0.0	34.0 / 40.5	28.0 / 35.2	-7.2

Annual all-type citations. Pre-fit RMSE describes the final fit using all pre-years; the held-out prediction uses only earlier years. Post is the calendar year after disclosure. The gap is actual minus synthetic citations. No sampling or randomization interval is implied. Full donor weights, paths and sensitivity estimates are in the repository.

The synthetic comparisons are mixed and do not resolve the counterfactual problem. Parental leave is close to its synthetic comparison and electrification is below it. The held-out pre-year predictions exceed actual citations in both cases. Close in-sample fit therefore does not establish a reliable untreated trajectory, and these gaps are not pooled into a causal estimate.